



From Zero to Hero: Conquering the Arm Neoverse

Brendan Bouffler (AWS)

Csaba Csoma (AWS)

John Linford (NVIDIA)

Matt Vaughn (AWS)

Conrad Hillairet (Arm Ltd)

Filippo Spiga (NVIDIA)



Compile and Execute

PART 1

I have no experience with Arm

Should I be worried ?

Overthinking:
**The art of creating problems that
weren't even there.**

**Most applications compile on Arm
architecture with little to no
modification**

- All major Linux distributions support Arm, with extensive library of common packages (AArch64).
- Upgrading to newer library / software versions might be needed.
- GNU Compiler Collection (GCC) and LLVM are fully supported.
- Commercial compilers for Arm are available.

GNU compilers are a solid option on Arm

- + GNU compilers are first class Arm compilers
 - Arm is one of the largest contributors to GCC
 - Focus on enablement and performance
 - Key for Arm to succeed in Cloud/Data center segment
- + GNU toolchain ships with ACFL
 - Good Fortran support
 - Bug fixes and performance improvements in upcoming GNU releases
- + Arm significant contributor to GNU



NVIDIA HPC Compiler Suite supporting Neoverse

- + Former PGI Compilers & Tools evolved into the NVIDIA HPC SDK™
- + Released multiple times year
 - Bundled with CPU and GPU libraries, Communication Library and Tools.
 - Free of charge and commercially supported.
- + Support all three major CPU families used in HPC systems
 - x86_64
 - aarch64 (sbsa)
 - ppc64le
- + Delivers state-of-the-art SIMD vectorization for both V1 and V2



arm COMPILER

Arm provided C/C++/Fortran compiler with best-in-class performance



Compilers tuned for Scientific Computing and HPC



Latest features and performance optimizations



Freely available

Tuned for Scientific Computing and HPC workloads

- Processor-specific optimizations for Neoverse-based platforms
- Optimal shared-memory parallelism using latest Arm-optimized OpenMP runtime

Linux user-space compiler with latest features

- C++ 17 and Fortran 2003 language support with OpenMP 4.5
- Support for Armv8-A and SVE architecture extension
- Based on LLVM and Flang, a leading open-source compiler project

Freely available on most Linux distributions

Documentation

- Fortran : <https://developer.arm.com/documentation/101380/2210>
- C/C++ : <https://developer.arm.com/documentation/101458/2210>

Compilers

How do I call them ?



gcc

g++

gfortran

arm
COMPILER
FOR LINUX

armclang

armclang++

armflang



nvc

nvc++

nvfortran

Edit your Makefile, CFLAGS and FCFLAGS

Compiler flags might be different between ifort, gfortran, armflang, nvfortran and frt

-qopenmp

-mp

-Kopenmp

-fopenmp



Ensure that you are using the correct compiler options.

No SMT. Relaxed memory model.

Compiler flags



FLAG	DESCRIPTION	
-mcpu=native or neoverse-v1, neoverse-v2, neoverse-512tvb	Target the detected or specific CPU (the default mcpu is more generic/slower)	
-O1, -O2, O3, -Ofast	Optimization level	
-fopenmp	Turns on OpenMP	
-f(no-) vectorize acfl	-f(no)tree- vectorize gcc	Enables vectorization
-f(no-)stack-arrays	Place all automatic arrays on stack memory (enabled by default with -Ofast).	
-f(no-)inline-functions	Enables function inlining.	
-f(no-)associative-math	Allows the re-association of FP operands	
-f(no-)lto	Enables Link Time Optimizations (LTO)	
-f(no-)unroll-loops	Enables loop unrolling	
-ffast-math	Enables aggressive, lossy FP optimizations.	
-ffp-contract=fast -ffp-contract=off	Enables FP expression contraction(i.e. FMAs)	



arm PERFORMANCE LIBRARIES

Optimized BLAS, LAPACK and FFT for HPC applications



Freely available



Best-in-class performance



Validated with
NAG test suite

Arm provided 64-bit Armv8-A math libraries

- Optimized BLAS, LAPACK, FFT and math.h routines
- FFTW compatible interface for FFT routines
- Sparse linear algebra and batched BLAS support
- Minimum Linux kernel version 5.10 is needed for SVE

Best-in-class serial and parallel performance

- Generic Armv8-A optimizations by Arm
- Tuned for Arm Neoverse family of processors

Validated by Arm Engineers

- Validated with NAG's test suite, a de-facto standard
- Community supported

Documentation

- <https://developer.arm.com/documentation/102620/0100>



Arm Performance Libraries

Optimized BLAS, LAPACK and FFT



FLAG	DESCRIPTION
<code>-I\$ARMPL_DIR/include</code>	Include directory
<code>-L\$ARMPL_DIR/lib -larmpl_mp</code>	Link against ArmPL Here OpenMP version (_mp)
<code>-L\$ARMPL_DIR/lib -lamath</code>	Link against libamath
<code>(-lgfortran) -lm</code>	Always finish by linking with -lm

arm
COMPILER
FOR LINUX

FLAG	DESCRIPTION
<code>-armpl</code>	Use with <code>-mcpu(=native)</code>
<code>-armpl=serial</code>	Single threaded version
<code>-armpl=parallel</code>	Default with <code>-fopenmp</code>



FLAG	DESCRIPTION
<code>-I\$ARMPL_DIR/include</code>	Include directory
<code>-L\$ARMPL_DIR/lib -larmpl_mp</code>	Link against ArmPL Here OpenMP version (_mp)
<code>-lm -lnvf -lnvc -lrt</code>	Always finish by linking with these libraries

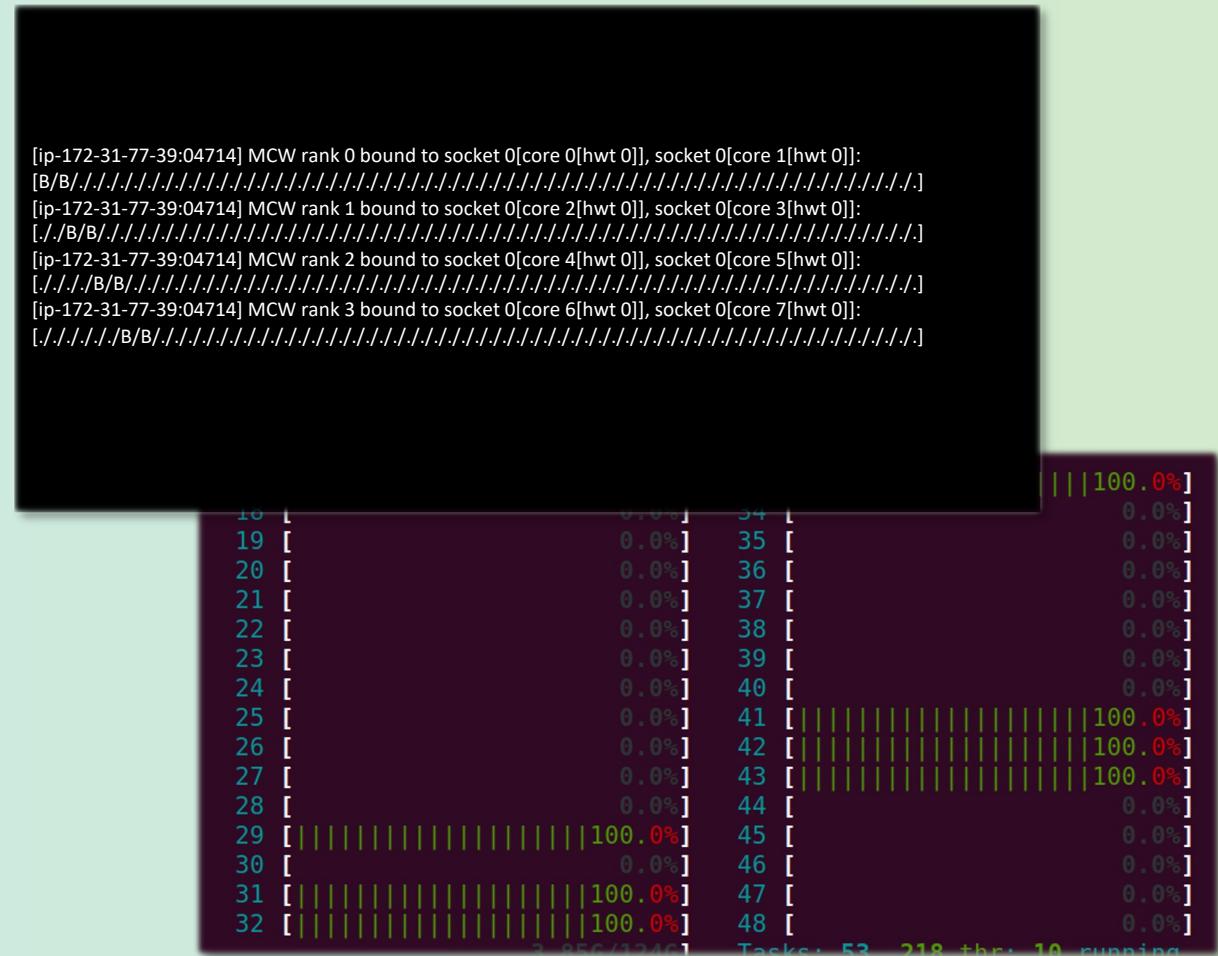
MPI & OpenMP binding and placement

On which cores am I running?

Some flags, commands and environment variables

- + --report-bindings
- + htop
- + OMP_DISPLAY_ENV
- + OMP_DISPLAY_AFFINITY

- + --map-by ppr:32:node:PE=2
- + OMP_NUM_THREADS=2
- + OMP_PROC_BIND=close
- + OMP_PLACES=cores



Useful OpenMP Runtime Environment Variables

OMP_NUM_THREADS

- Normally defaults to the number of cores on the system, i.e. 64 on HPC7g/C7g.16xlarge

OMP_DISPLAY_ENV

- Set to “true” or “verbose” to see values of run-time variables

OMP_PROC_BIND

- Strongly recommended to set OMP_PROC_BIND to “true”, “close” or “spread” depending on needs
- If unset, the threads are not pinned to cores and may migrate around the system

OMP_PLACES

- Description of unit of pinning ‘threads’, ‘cores’ or ‘sockets’

OMP_DYNAMIC=true and OMP_NESTED=true

- Necessary for “nested parallelism” (multi-threaded functions from within an already parallel environment)

Guided Hands-Ons || Bring Your Own Code

50 MINUTES

Part 1 – Practical hands-on

Goal: compile the code(s), check correctness, experiment with different compilers and libraries (*which combination is faster?*), single/multi-instances

Options:

1. “BYOC” - Bring Your Own Code
2. code_saturne play-through
3. Other popular HPC apps (follow the links)



<https://github.com/aws/aws-graviton-getting-started/tree/main/HPC>



<https://github.com/arm-hpc-devkit/nvidia-arm-hpc-devkit-users-guide>

Code_Saturne

Reference - Build instructions on x86 with Intel toolchain

Install

```
source ~/intel/oneapi/setvars.sh
CC="mpiicc"
CXX="mpiicpc"
FC="ifort"
DEST=$HOME/code_saturne_7.2.0
mkdir build
cd build
../configure CC="$CC" CXX="$CXX" FC="$FC" \
    --with-blas=$MKLROOT --prefix=$DEST \
    --disable-gui --without-med \
    --without-hdf5 --without-cgns \
    --without-metis --disable-salome \
    --without-salome --without-eos \
    --disable-static --enable-long-gnu
make
make install
```

Run

- Download test cases <https://github.com/code-saturne/saturne-open-cases>
- Two test cases:
 - C016 allows very fast computations (single core)
 - F128 more precise simulations (single node)
- **First step:** generate the meshes that will be used later
 - cd \$REPBASE/BENCH_C016_PREPROCESS/DATA
 - \$DEST/bin/code_saturne run
 - Outputted data are in \$REPBASE/BENCH_C016_PREPROCESS/RESU/extrude_16
 - C016 : has 1 024 cells as an input and 17 408 as the output
 - cd \$REPBASE/BENCH_F128_PREPROCESS/DATA
 - \$DEST/bin/code_saturne run
 - Outputted data are in \$REPBASE/BENCH_F128_PREPROCESS/RESU/extrude_128
 - F128 : 100 040 cells in input, 12 905 160 cells in output
- **Second step:** run C016 on one node with 48 cores and 1 OMP thread:
 - cd \$REPBASE/BENCH_C016_01/DATA
 - \$DEST/bin/code_saturne submit --wckey dtsi:undefined --nodes=1 -n 48 --exclusive --ntasks-per-node=48 --cpus-per-task=1



Adapt configure command to your environment

x86 to aarch64

```
CC="mpiicc"
CXX="mpiicpc"
FC="ifort"
DEST=$HOME/code_saturne_7.2.0
mkdir build
cd build
./configure CC="$CC" CXX="$CXX" FC="$FC" \
    --with-blas=$MKLROOT --prefix=$DEST \
    --disable-gui --without-med \
    --without-hdf5 --without-cgns \
    --without-metis --disable-Salome \
    --without-salome --without-eos \
    --disable-static --enable-long-gnu
make
make install
```

```
CC=mpicc CXX=mpicxx FC=armflang ./configure \
--with-blas=$ARMPL_DIR --prefix=$PWD/build \
--disable-gui --without-med \
--without-hdf5 --without-cgns \
--without-metis --disable-Salome \
--without-salome --without-eos \
--disable-static --enable-long-gnu
```

Adapt configure command to your environment

x86 to aarch64

```
[...]
checking for python3 extension module
directory...
${exec_prefix}/lib64/python3.7/site-packages
checking for dlopen... dlopen
checking for MKL libraries... no
checking for threaded ATLAS BLAS... no
checking for ATLAS BLAS... no
checking for legacy C BLAS... no
configure: error: in `/home/ec2-
user/code_saturne-7.2.0':
configure: error: BLAS support is requested,
but test for BLAS failed!
See `config.log' for more details
```

Need to introduce support for
ArmPL

Modify one file

Adapt ATLAS and MKL detection
mechanisms

Adapt configure command to your environment

x86 to aarch64

```
[...]
checking for python3 extension module directory...
${exec_prefix}/lib64/python3.7/site-packages
checking for dlopen... dlopen
checking for MKL libraries... no
checking for threaded ATLAS BLAS... no
checking for ATLAS BLAS... no
checking for legacy C BLAS... no
configure: error: in `/home/ec2-user/code_saturne-
7.2.0':
configure: error: BLAS support is requested, but test
for BLAS failed!
See `config.log' for more details
```

Need to introduce support for
ArmPL

Modify one file

Adapt ATLAS and MKL detection
mechanisms

Adapt configure command to your environment

x86 to aarch64

```
[...]  
${exec_prefix}/lib64/python3.7/site-packages  
checking for dlopen... dlopen  
checking for MKL libraries... no  
checking for threaded ATLAS BLAS... no  
checking for ATLAS BLAS... no  
checking for legacy C BLAS... no  
checking for ArmPL libraries... yes  
checking for MPI (MPI compiler wrapper test)... yes  
checking for MPI I/O... yes  
checking for MPI in place... yes  
checking for MPI nonblocking barrier... yes  
CCMIO headers not found  
[...]
```

+—————
**Need to introduce support for
ArmPL**

+—————
Modify one file

+—————
**Adapt ATLAS and MKL detection
mechanisms**

Adapt configure command to your environment

x86 to aarch64

```
[...]
checking for a sed that does not truncate output...
/usr/bin/sed
checking for ar... ar
checking the archiver (ar) interface... ar
configure: sourcing config/cs_auto_flags.sh (test for
known compilers)
compiler 'mpicc' is NVIDIA compiler
compiler 'mpicxx' is NVIDIA compiler
compiler 'armflang' is NVIDIA compiler
checking how to print strings... printf
[...]
```

Bug in compiler detection mechanism

Arm Compiler for Linux detection overwritten by NVIDIA detection

Two files to modify (~10-20 lines)

Adapt configure command to your environment

x86 to aarch64

```
[...]
checking for ar... ar
checking the archiver (ar) interface... ar
configure: sourcing config/cs_auto_flags.sh
(test for known compilers)
compiler 'mpicc' is Arm C compiler
compiler 'mpicxx' is Arm C++
compiler 'armflang' is Arm Fortran compiler
checking how to print strings... printf
[...]
```

Bug in compiler detection mechanism

Arm Compiler for Linux detection overwritten by NVIDIA detection

Two files to modify (~10-20 lines)

Compile

x86 to aarch64

```
make -j 4 > resComp >&1  
make install  
  
grep -i err ./resComp | wc -l  
0
```

No problem, compiles OoB

Flat MPI version

x86 to aarch64

```
CC=mpicc CXX=mpicxx FC=gfortran ./configure --  
prefix=$PWD/build-noomp --disable-gui --without-hdf5 --  
without-cgns --without-med --without-metis --disable-static  
--disable-salome --without-salome --without-eos --enable-long-  
gnum --with-blas=$ARMPL_DIR --disable-openmp > resConf 2>&1  
  
grep -i openmp ./resConf  
[...]  
AArch64_RHEL-7_aarch64-linux/bin/./lib/gcc/aarch64-linux-  
gnu/11.2.0/.../... -lgfortran -lm  
checking for OpenMP (C)... yes  
checking for OpenMP (Fortran)... yes  
checking for Fortran libraries of gfortran... (cached) OpenMP  
support: no  
OpenMP support: yes  
OpenMP accelerator support: no  
OpenMP Fortran support: yes  
Auto load environment modules: binutils/11.2.0 gnu/11.2.0  
armpl/22.1.0 openmpi/gcc/4.1.4
```

Bug in OpenMP configuration
not Arm specific
also happens on x86

One file to modify (~1 line)

Flat MPI version

x86 to aarch64

```
CC=mpicc CXX=mpicxx FC=gfortran ./configure --  
prefix=$PWD/build-noomp --disable-gui --without-hdf5 --  
without-cgns --without-med --without-metis --disable-static -  
--disable-salome --without-salome --without-eos --enable-long-  
gnum --with-blas=$ARMPL_DIR --disable-openmp > resConf 2>&1  
  
grep -i openmp ./resConf  
[...]  
OpenMP support: no  
OpenMP support: no  
Auto load environment modules: binutils/11.2.0 gnu/11.2.0  
armpl/22.1.0 openmpi/gcc/4.1.4  
  
make > resComp 2>&1  
grep fopenmp ./resComp | wc -l  
0
```

Bug in OpenMP configuration
not Arm specific
also happens on x86

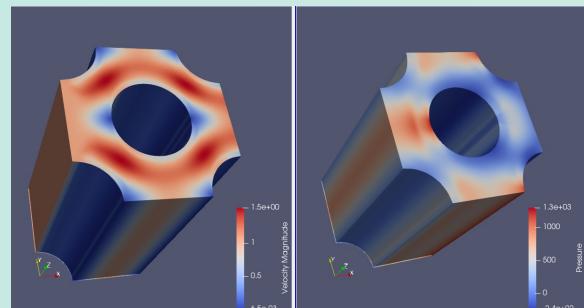
One file to modify (~1 line)

Run x86 to aarch64

First step : generate the meshes that will be used later
cd \$REPBASE/BENCH_F128_PREPROCESS/DATA
\$DEST/bin/code_saturne run

Second step: run testcase C016 on one node with 48 cores
and 1 thread per MPI task :
cd \$REPBASE/BENCH_C016_01/DATA

\$DEST/bin/code_saturne submit --wckey dtsi:undefined --nodes=1 -n 48 --exclusive --ntasks-per-node=48 --cpus-per-task=1

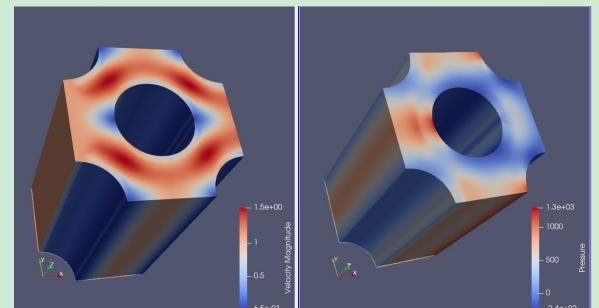


No problem at runtime

First step : generate the meshes that will be used later
cd \$REPBASE/BENCH_F128_PREPROCESS/DATA
\$DEST/bin/code_saturne run

Second step: run testcase C016 on one node with 64 cores
and 1 thread per MPI task :

cd \$REPBASE/BENCH_C016_01/DATA
\$DEST/bin/code_saturne submit -n 64 --cpus-per-task=1



Another run with MPI & OpenMP

x86 to aarch64

Not Arm specific

```
code_saturne submit -n 4 --nt 2 --cpus-per-task=2
```

PID	USER	PRI	NI	VIRT	RES	SHR-S	CPU%	MEM%	TIME+	Command
7721	ec2-user	20	0	643M	159M	23476	R	200.	0.1	/cs solver --mpi
7720	ec2-user	20	0	643M	159M	23264	R	200.	0.1	0:05.84 ./cs solver --mpi
7718	ec2-user	20	0	670M	186M	23160	R	200.	0.1	0:05.83 ./cs solver --mpi

